

# Rigorous Explanations for Machine Learning Models

**Joao Marques-Silva**

(joint work with A. Ignatiev and N. Narodytska)

University of Lisbon, Portugal

AITP 2019 Conference

Obergurgl, Austria

April 2019

# Progress in automated reasoning

- Automated reasoners (AR):
  - SAT
  - ILP

# Progress in automated reasoning

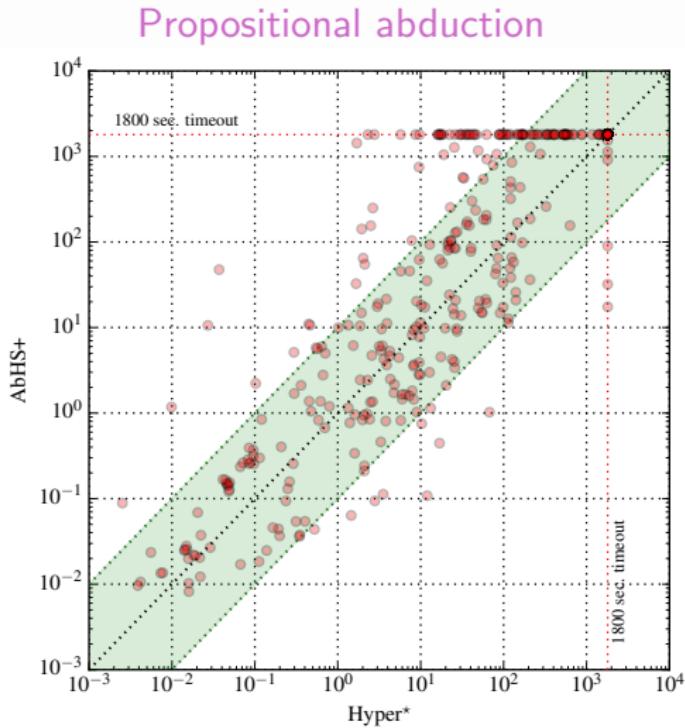
- Automated reasoners (AR):
  - SAT
  - ILP
  - ASP
  - SMT
  - FOL

# Progress in automated reasoning

- Automated reasoners (AR):
  - SAT
  - ILP
  - ASP
  - SMT
  - FOL
  - Reasoners as oracles
  - Reasoners within reasoners

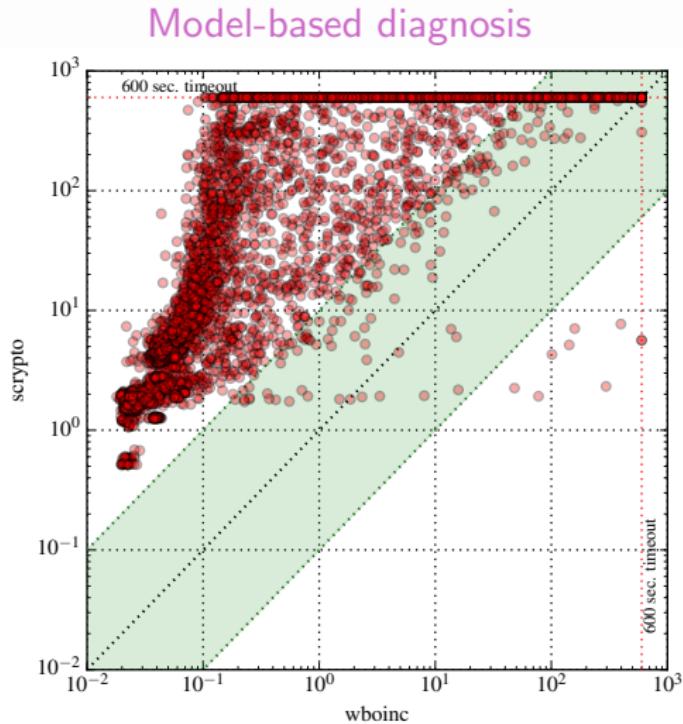
# Progress in automated reasoning & our work

- Automated reasoners (AR):
  - SAT
  - ILP
  - ASP
  - SMT
  - FOL
  - Reasoners as oracles
  - Reasoners within reasoners



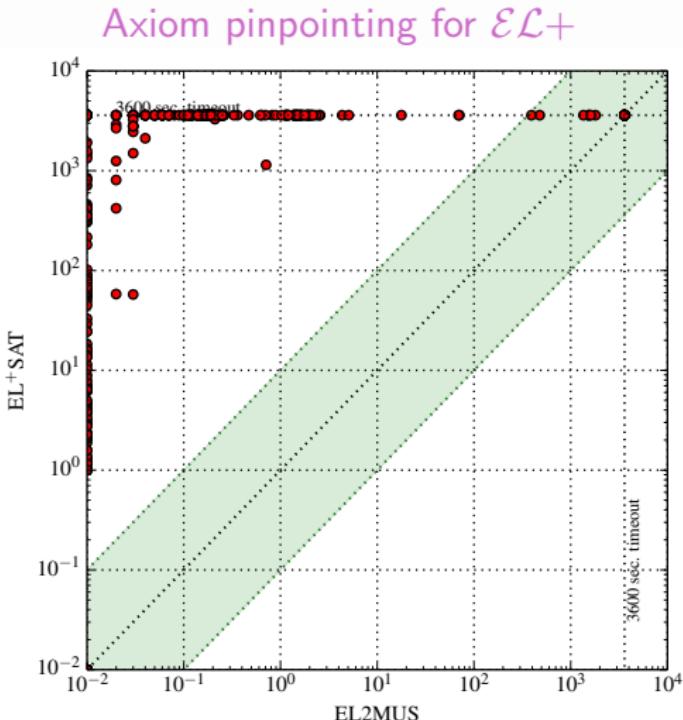
# Progress in automated reasoning & our work

- Automated reasoners (AR):
  - SAT
  - ILP
  - ASP
  - SMT
  - FOL
  - Reasoners as oracles
  - Reasoners within reasoners



# Progress in automated reasoning & our work

- Automated reasoners (AR):
  - SAT
  - ILP
  - ASP
  - SMT
  - FOL
  - Reasoners as oracles
  - Reasoners within reasoners



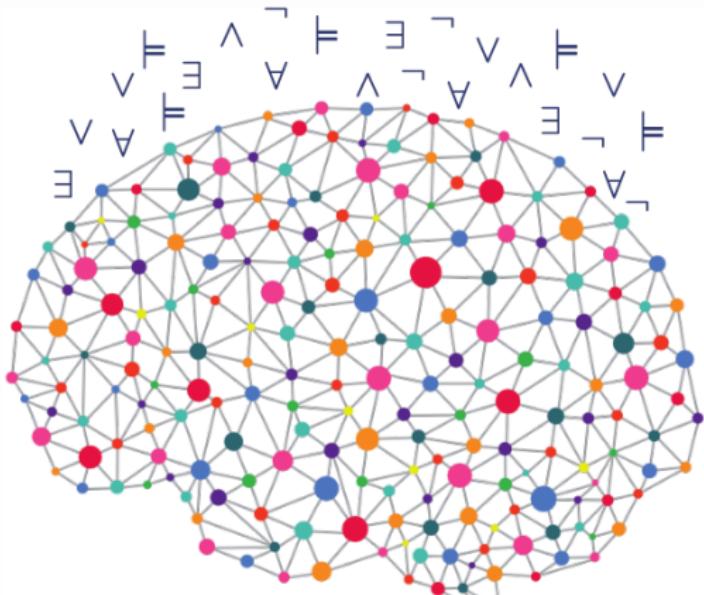
# The question: how can AR improve ML's robustness?

Moshe Vardi

## Machine learning and logic: Fast and slow thinking

**ABSTRACT.** There is a recent perception that computer science is undergoing a Kuhnian paradigm shift, with CS as a model-driven science being replaced as a data-driven science. I will argue that, in general new scientific theories refine old scientific theories, rather than replace them. Thus, data-driven CS and model-driven CS complement each other, just as fast thinking and slow thinking complement each other in human thinking, as explicated by Daniel Kahneman. I will then use automated vehicles as an example, where in recent years, car makers and tech companies have been racing to be the first to market. In this rush there has been little discussion of how to obtain scalable standardization of safety assurance, without which this technology will never be commercially deployable. Such assurance requires formal methods, and combining machine learning with logic is the challenge of the day.

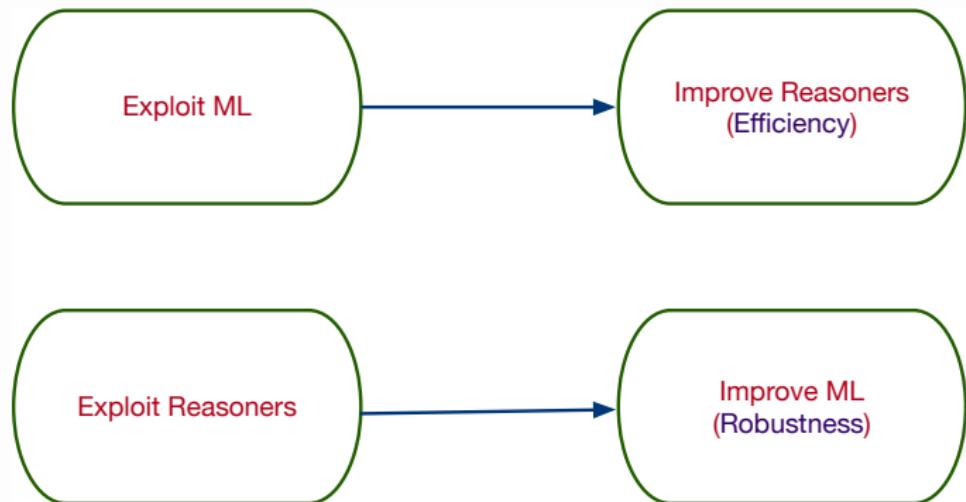
M. Vardi, MLMFM'18 Summit



# Machine learning vs. automated reasoning



# Machine learning vs. automated reasoning



## Our work ...

- Focus on **classification** problems

## Our work ...

- Focus on **classification** problems
- **Globally correct** (ie **rigorous**) **explanations** for predictions made

## Our work ...

- Focus on **classification** problems
- **Globally correct** (ie **rigorous**) **explanations** for predictions made
- **Disclaimer:** first inroads into ML & XAI;  
comments welcome

# Outline

Successes & Pitfalls of ML

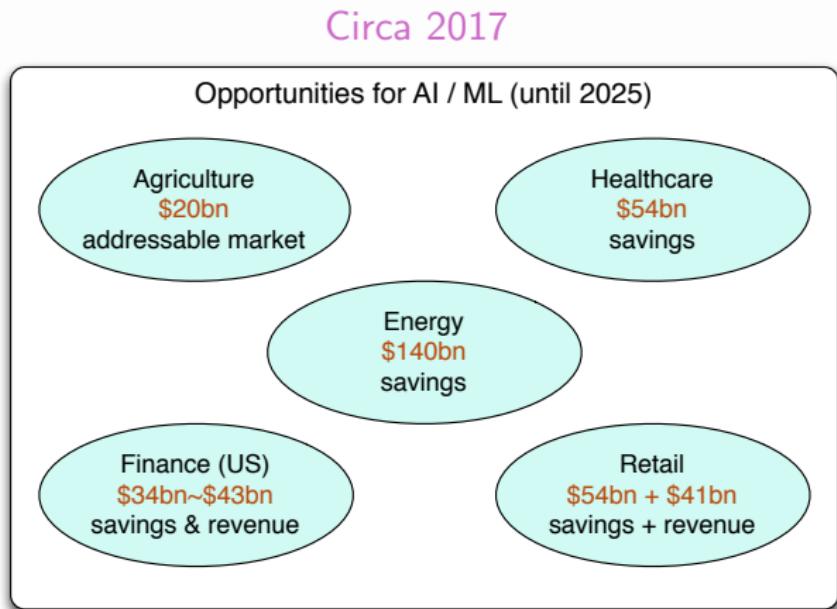
Explainable AI

Explanations with Abductive Reasoning

Results

# Some ML successes & expectations

- IBM Watson
- Deepmind AlphaGo
  - & AlphaZero
- Image Recognition
- Speech Recognition
- Financial Services
- Medical Diagnosis
- ...



## Many more applications expected



source: Google

## Many more applications expected



source: Google



source: Wikipedia



©DARPA

But ML models are **brittle**



Eykholt et al'18



Aung et al'17

## But ML models are **brittle**



Eykholt et al'18

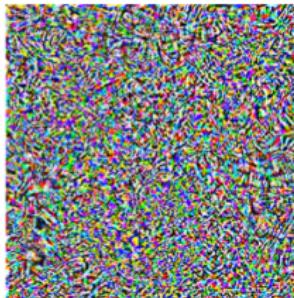


Aung et al'17



“pig”

+ 0.005 ×



=



“airliner”

Source: [http://gradientscience.org/intro\\_adversarial/](http://gradientscience.org/intro_adversarial/)

Also, some ML models are **interpretable**

decision|rule lists|sets  
decision trees

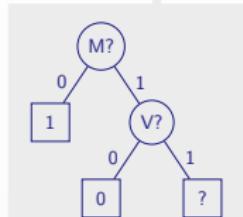
Ex.	Vacation (V)	Concert (C)	Meeting (M)	Expo (E)	Hike (H)
$e_1$	0	0	1	0	0
$e_2$	1	0	0	0	1
$e_3$	0	0	1	1	0
$e_4$	1	0	0	1	1
$e_5$	0	1	1	0	0
$e_6$	0	1	1	1	0
$e_7$	1	1	0	1	1

# Also, some ML models are **interpretable**

decision|rule lists|sets  
decision trees

if  $\neg$ Meeting then Hike  
if  $\neg$ Vacation then  $\neg$ Hike

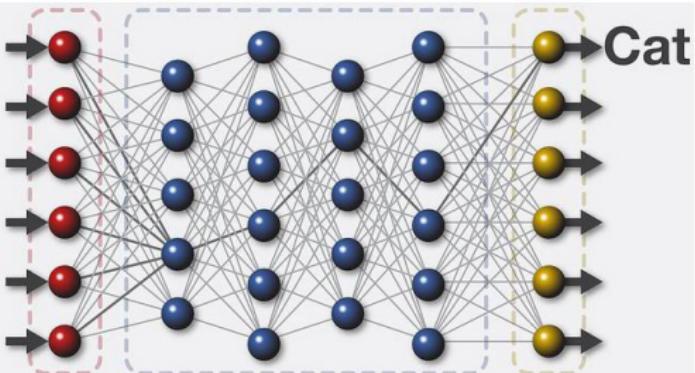
Ex.	Vacation (V)	Concert (C)	Meeting (M)	Expo (E)	Hike (H)
$e_1$	0	0	1	0	0
$e_2$	1	0	0	0	1
$e_3$	0	0	1	1	0
$e_4$	1	0	0	1	1
$e_5$	0	1	1	0	0
$e_6$	0	1	1	1	0
$e_7$	1	1	0	1	1



But other ML models are **not** (interpretable)...



©DARPA



Why does the NN predict a cat?

## Sample of ongoing efforts

- **Verification of NNs:**
  - Sound vs. unsound vs. complete [M.P. Kumar, VMCAI'19]
  - E.g. Reluplex: dedicated reasoning within SMT solver
- **Explanations for non-interpretable (ie black-box) models:**
  - Until recently, most approaches heuristic-based

# Outline

Successes & Pitfalls of ML

Explainable AI

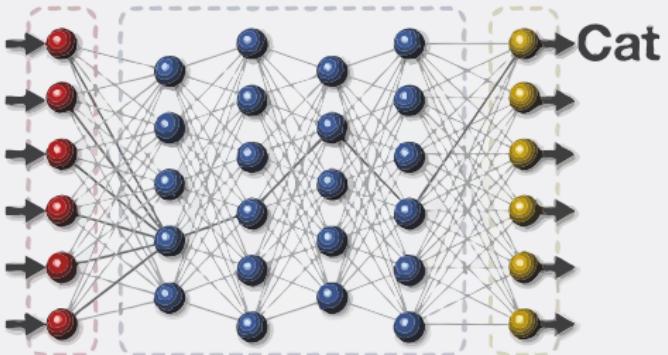
Explanations with Abductive Reasoning

Results

# What is eXplainable AI (XAI)?

# What is eXplainable AI (XAI)?

## Machine Learning System



This is a cat.

Current Explanation

This is a cat:

- It has fur, whiskers, and claws.
- It has this feature:



XAI Explanation

# Why XAI?

REGULATION (EU) 2016/679 OF THE EUROPEAN PARLIAMENT AND OF THE COUNCIL

of 27 April 2016

on the protection of natural persons with regard to the processing of personal data and on the free movement of such data, and repealing Directive 95/46/EC (General Data Protection Regulation)

(Text with EEA relevance)

# Why XAI?

REGULATION (EU) 2016/679 OF THE EUROPEAN PARLIAMENT AND OF THE COUNCIL

of 27 April 2016

on the protection of natural persons with regard to the processing of personal data and on the free movement of such data, and repealing Directive 95/46/EC (General Data Protection Regulation)

(Text with EEA relevance)

European Union regulations on algorithmic decision-making  
and a “right to explanation”

Bryce Goodman,<sup>1\*</sup> Seth Flaxman,<sup>2</sup>

# Why XAI?

REGULATION (EU) 2016/679 OF THE EUROPEAN PARLIAMENT AND OF THE COUNCIL

of 27 April 2016

on the protection of natural persons with regard to the processing of personal data and on the free movement of such data, and repealing Directive 95/46/EC (General Data Protection Regulation)

(Text with EEA relevance)

European Union regulations on algorithmic decision-making  
and a “right to explanation”

Bryce Goodman,<sup>1\*</sup> Seth Flaxman,<sup>2</sup>

■ We summarize the potential impact that the European Union’s new General Data Protection Regulation will have on the routine use of machine-learning algorithms. Slated to take effect as law across the European Union in 2018, it will place restrictions on automated individual decision making (that is, algorithms that make decisions based on user-level predictors) that “significantly affect” users. When put into practice, the law may also effectively create a right to explanation, whereby a user can ask for an explanation of an algorithmic decision that significantly affects them. We argue that while this law may pose large challenges for industry, it highlights opportunities for computer scientists to take the lead in designing algorithms and evaluation frameworks that avoid discrimination and enable explanation.

# Why XAI?

REGULATION (EU) 2016/679 OF THE EUROPEAN PARLIAMENT AND OF THE COUNCIL

of 27 April 2016

on the protection of natural persons with regard to the processing of personal data and on the free movement of such data, and repealing Directive 95/46/EC (General Data Protection Regulation)

(Text with EEA relevance)

European Union regulations on algorithmic decision-making  
and a “right to explanation”

Bryce Goodman,<sup>1\*</sup> Seth Flaxman,<sup>2</sup>

## SUMMIT ON MACHINE LEARNING MEETS FORMAL METHODS

■ We summarize the potential impact that the European Union’s new General Data Protection Regulation will have on the routine use of machine-learning algorithms. Slated to take effect as law across the European Union in 2018, it will place restrictions on automated individual decision making (that is, algorithms that make decisions based on user-level predictors) that “significantly affect” users. When put into practice, the law may also effectively create a right to explanation, whereby a user can ask for an explanation of an algorithmic decision that significantly affects them. We argue that while this law may pose large challenges for industry, it highlights opportunities for computer scientists to take the lead in designing algorithms and evaluation frameworks that avoid discrimination and enable explanation.

# Why XAI?

REGULATION (EU) 2016/679 OF THE EUROPEAN PARLIAMENT AND OF THE COUNCIL

of 27 April 2016

on the protection of natural persons with regard to the processing of personal data and on the free movement of such data, and repealing Directive 95/46/EC (General Data Protection Regulation)

(Text with EEA relevance)

European Union regulations on algorithmic decision-making  
and a “right to explanation”

Bryce Goodman,<sup>1\*</sup> Seth Flaxman,<sup>2</sup>

SUMMIT ON MACHINE LEARNING  
MEETS FORMAL METHODS

■ We summarize the potential impact that the European Union's new General Data Protection Regulation will have on the routine use of machine-learning algorithms. Slated to take effect as law across the European Union in 2018, it will place restrictions on automated individual decision making (that is, algorithms that make decisions based on user-level predictors) that "significantly affect" users. When put into practice, the law may also effectively create a right to explanation, whereby a user can ask for an explanation of an algorithmic decision that significantly affects them. We argue that while this law may pose large challenges for industry, it highlights opportunities for computer scientists to take the lead in designing algorithms and evaluation frameworks that avoid discrimination and enable explanation.

## Explainable Artificial Intelligence (XAI)



David Gunning  
DARPA/I2O  
Program Update November 2017



# Relevancy of XAI



## MIT Technology Review

The Dark Secret at the Heart of AI  
Will Knight  
April 11, 2017

Intelligent Machines Are Asked to Explain How Their Minds Work



Richard Waters  
July 11, 2017

## Entrepreneur

Elon Musk and Mark Zuckerberg Are Arguing About AI -- But They're Both Missing the Point

Artur Kuiulan  
July 28, 2017



DARPA's XAI seeks explanations from autonomous systems

Geoff Fein  
November 16, 2017

COMPUTERWORLD  
Oracle quietly researching 'Explainable AI'

George Nott  
May 5, 2017



## Inside DARPA's Push to Make Artificial Intelligence Explain Itself

Sara Castellanos and Steven Norton  
August 10, 2017

## The Register

You better explain yourself, mister:  
DARPA's mission to make an accountable AI

Dan Robinson  
September 29, 2017



Team investigates artificial intelligence, machine learning in DARPA project  
Lisa Daigle  
June 14, 2017

## Military EMBEDDED SYSTEMS

NOVATEK  
Ghosts in the Machine  
Christina Couch  
October 25, 2017



## SCIENTIFIC AMERICAN

Demystifying the Black Box That Is AI  
Ariel Bleicher  
August 9, 2017

## The New York Times Magazine



Can A.I. Be Taught to Explain Itself?  
Cliff Kuang  
November 21, 2017



## ExecutiveBiz

Charles River Analytics-Led Team Gets DARPA Contract to Support Artificial Intelligence Program  
Ramona Adams  
June 13, 2017

## FAST COMPANY

Why The Military And Corporate America Want To Make AI Explain Itself  
Steven Melendez  
June 22, 2017



How AI detectives are cracking open the black box of deep learning

Paul Voosen  
July 6, 2017



# Relevancy of XAI & hundreds(?) of recent papers



**MIT  
Technology  
Review**

**The Dark Secret at  
the Heart of AI**  
Will Knight  
April 11, 2017

Intelligent Machines  
Are Asked to Explain  
How Their Minds  
Work

Richard Waters  
July 11, 2017

**FT**



INANCIA

**Entrepreneur**

Elon Musk and Mark  
Zuckerberg Are Arguing  
About AI -- But They're  
Both Missing the Point

Artur Kuiulan  
July 28, 2017



Jane's

DARPA's XAI seeks  
explanations from  
autonomous systems  
Geoff Fein  
November 16, 2017

**COMPUTERWORLD**

Oracle quietly  
researching  
'Explainable AI'  
George Nott  
May 5, 2017



THE WALL STREET JOURNAL  
**WSJ**

**Inside DARPA's Push to  
Make Artificial Intelligence  
Explain Itself**  
Sara Castellanos and Steven  
Norton  
August 10, 2017

**The Register**

You better explain  
yourself, mister:  
DARPA's mission to  
make an accountable AI  
Dan Robinson  
September 29, 2017



**The New York Times Magazine**



**Can A.I. Be  
Taught to  
Explain Itself?**  
Cliff Kuang  
November 21, 2017

**ExecutiveBiz**

Charles River Analytics-Led  
Team Gets DARPA Contract to  
Support Artificial Intelligence  
Program  
Ramona Adams  
June 13, 2017



**Military**  
EMBEDDED SYSTEMS

Team investigates artificial  
intelligence, machine learning  
in DARPA project  
Lisa Daigle  
June 14, 2017

**NOVATEXT**

Ghosts in the Machine  
Christina Couch  
October 25, 2017

**FAST COMPANY**

Why The Military And  
Corporate America Want  
To Make AI Explain  
Itself  
Steven Melendez  
June 22, 2017



**SCIENTIFIC  
AMERICAN**

Demystifying the  
Black Box That Is AI  
Ariel Bleicher  
August 9, 2017



How AI detectives are cracking open  
the black box of deep learning

Paul Voosen  
July 6, 2017

**Science**  
AAAS

# How to XAI?

Main challenge: **black-box models**

Heuristic approaches, e.g. **LIME & Anchor**

[Guerreiro et al., KDD'16, AAAI'18]

- Compute **local** explanations ...

# How to XAI?

Main challenge: **black-box models**

Heuristic approaches, e.g. **LIME & Anchor**

[Guerreiro et al., KDD'16, AAAI'18]

- Compute **local** explanations ...
- ... offer **no** guarantees

# How to XAI?

Main challenge: **black-box models**

Heuristic approaches, e.g. **LIME & Anchor**

[Guerreiro et al., KDD'16, AAAI'18]

- Compute **local** explanations ...
- ... offer **no** guarantees

Recent efforts on **rigorous** approaches

- **Compilation**-based, e.g. for BNCs

[Shih, Choi & Darwiche, IJCAI'18]

- ▶ Issues with scalability

- **Abduction**-based, e.g. for NNs

[Ignatiev, Narodytska, M.-S., AAAI'19]

- ▶ Issues with scalability

# How to XAI?

Main challenge: **black-box models**

Heuristic approaches, e.g. **LIME & Anchor**

[Guerreiro et al., KDD'16, AAAI'18]

- Compute **local** explanations ...
- ... offer **no** guarantees

Recent efforts on **rigorous** approaches

- **Compilation**-based, e.g. for BNCs

[Shih, Choi & Darwiche, IJCAI'18]

- ▶ Issues with scalability

- **Abduction**-based, e.g. for NNs

[Ignatiev, Narodytska, M.-S., AAAI'19]

- ▶ Issues with scalability, **but less significant**

## Some current challenges

- For heuristic methods: lack of rigor (more later)

## Some current challenges

- For heuristic methods: lack of rigor (more later)
- For rigorous methods: scalability, scalability, scalability...

# Outline

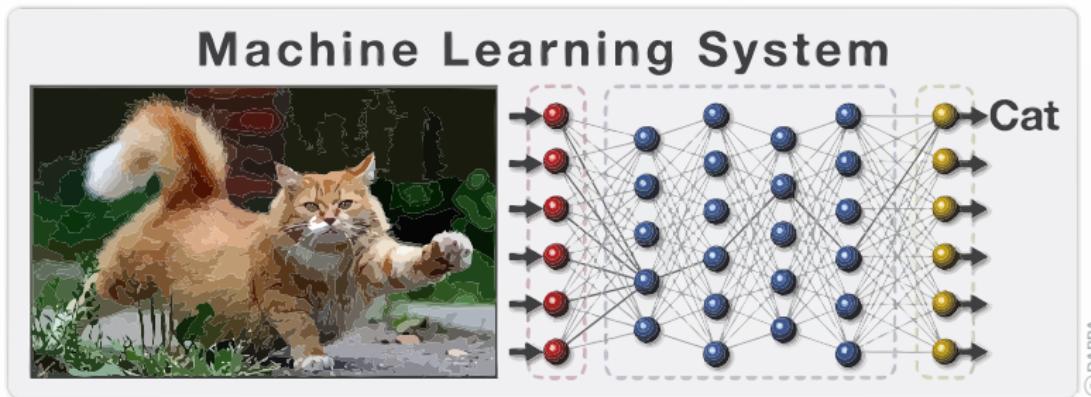
Successes & Pitfalls of ML

Explainable AI

Explanations with Abductive Reasoning

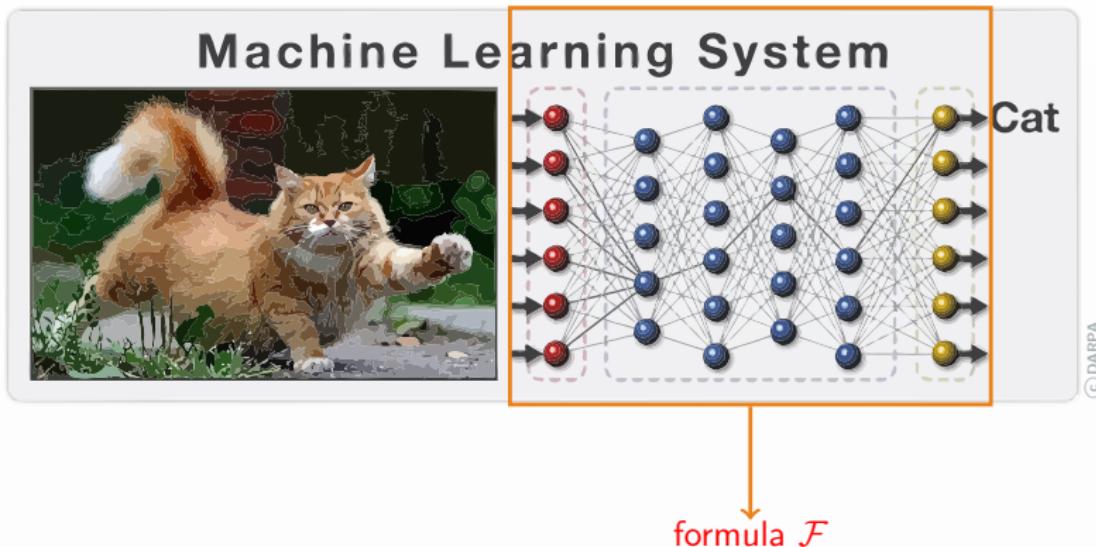
Results

# From ML model to logic

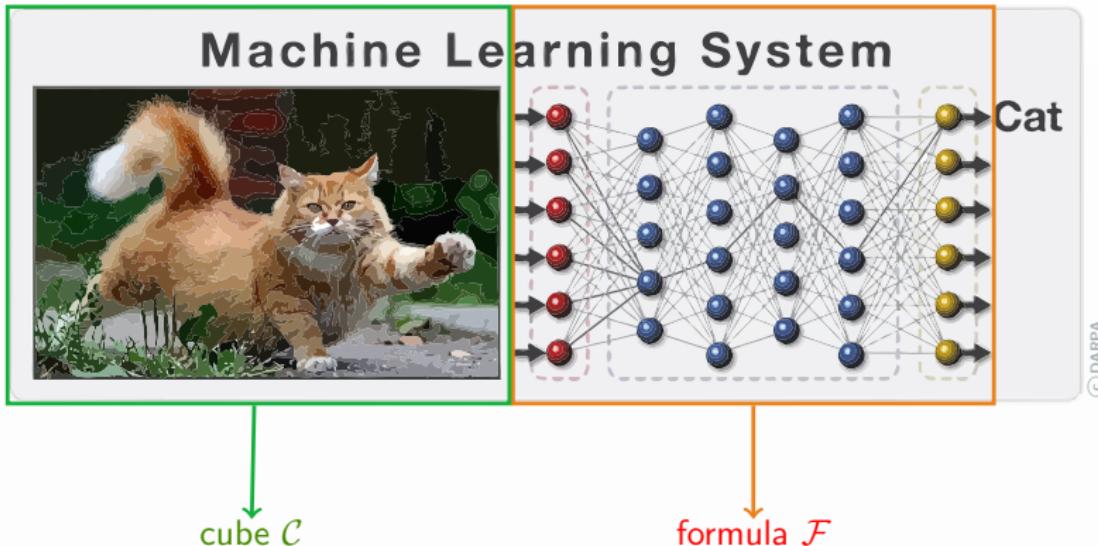


© DARPA

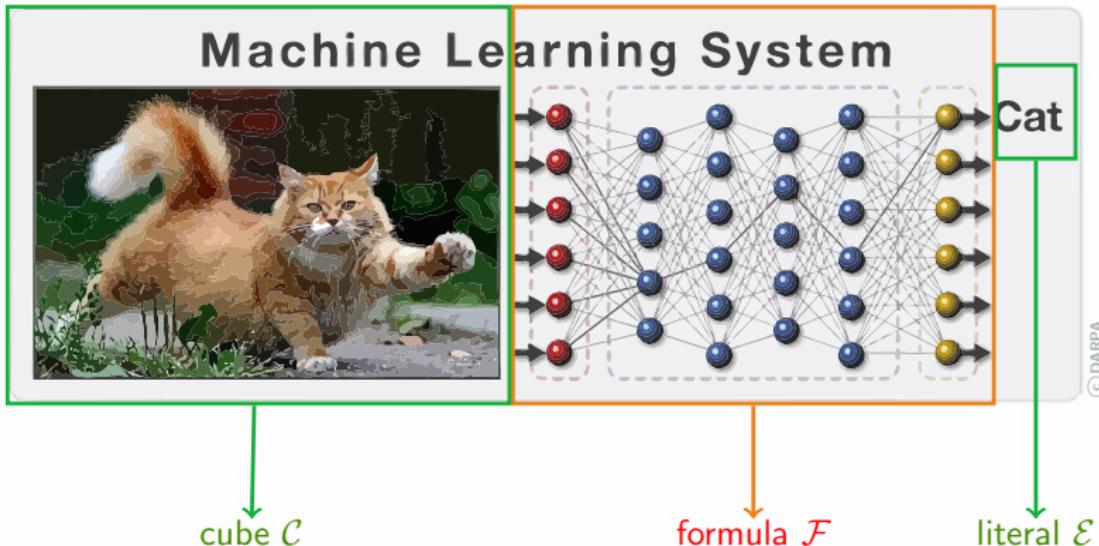
# From ML model to logic



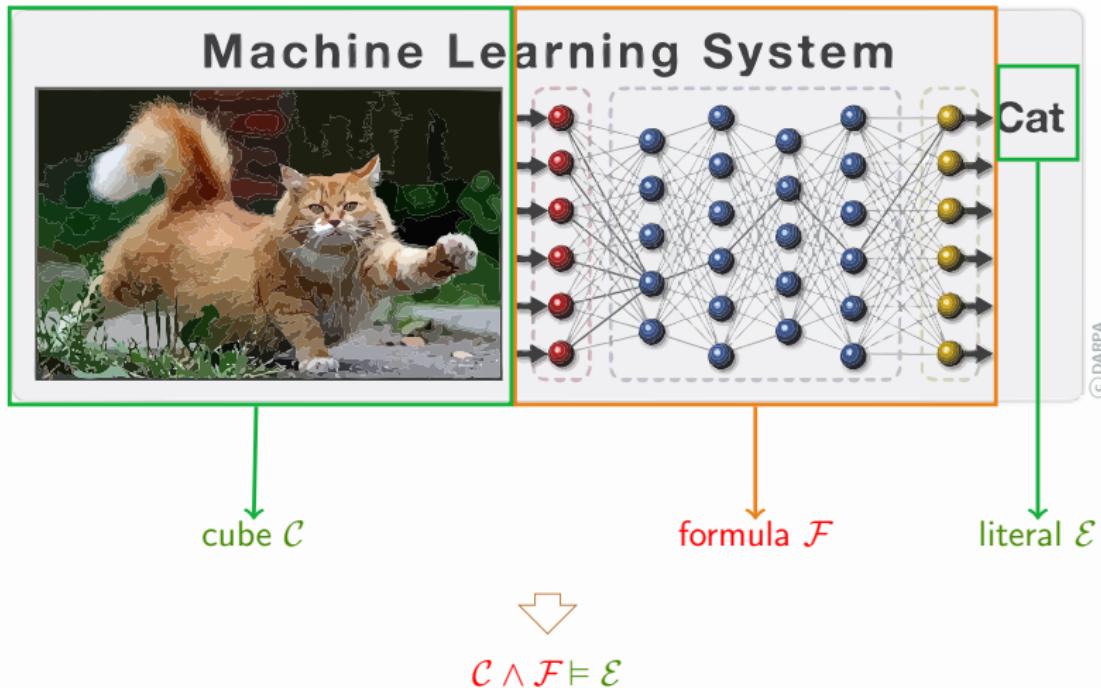
# From ML model to logic



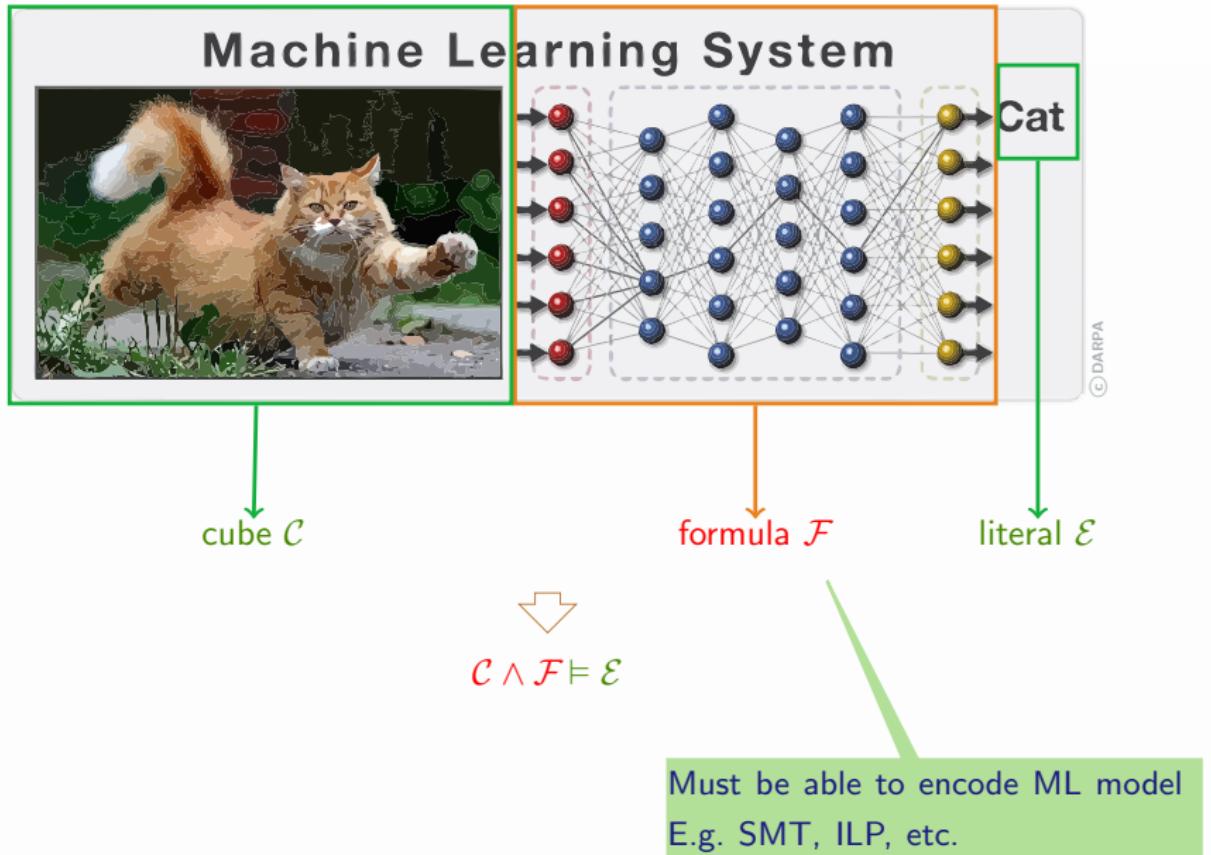
# From ML model to logic



# From ML model to logic



# From ML model to logic



# Abductive explanations of ML models

given a **classifier**  $\mathcal{F}$ , a **cube**  $\mathcal{C}$  and a **prediction**  $\mathcal{E}$ ,

## Abductive explanations of ML models

given a **classifier**  $\mathcal{F}$ , a **cube**  $\mathcal{C}$  and a **prediction**  $\mathcal{E}$ ,  
compute a (**subset-** or **cardinality-**) minimal  $\mathcal{C}_m \subseteq \mathcal{C}$  s.t.

# Abductive explanations of ML models

given a **classifier**  $\mathcal{F}$ , a **cube**  $\mathcal{C}$  and a **prediction**  $\mathcal{E}$ ,  
compute a (**subset-** or **cardinality-**) minimal  $\mathcal{C}_m \subseteq \mathcal{C}$  s.t.

$$\mathcal{C}_m \wedge \mathcal{F} \not\models \perp$$

and

$$\mathcal{C}_m \wedge \mathcal{F} \vDash \mathcal{E}$$

# Abductive explanations of ML models

given a **classifier**  $\mathcal{F}$ , a **cube**  $\mathcal{C}$  and a **prediction**  $\mathcal{E}$ ,  
compute a (**subset-** or **cardinality-**) minimal  $\mathcal{C}_m \subseteq \mathcal{C}$  s.t.

$$\mathcal{C}_m \wedge \mathcal{F} \not\models \perp$$

and

$$\mathcal{C}_m \wedge \mathcal{F} \vDash \mathcal{E}$$



**iterative explanation procedure**

# Computing primes

1.  $\mathcal{C}_m \wedge \mathcal{F} \not\models \perp$

# Computing primes

1.  $\mathcal{C}_m \wedge \mathcal{F} \not\models \perp$  — **tautology**

# Computing primes

1.  $\mathcal{C}_m \wedge \mathcal{F} \not\models \perp$  — **tautology**
2.  $\mathcal{C}_m \wedge \mathcal{F} \models \mathcal{E}$

# Computing primes

1.  $\mathcal{C}_m \wedge \mathcal{F} \not\models \perp$  — **tautology**
2.  $\mathcal{C}_m \wedge \mathcal{F} \models \mathcal{E} \Leftrightarrow \mathcal{C}_m \models (\mathcal{F} \rightarrow \mathcal{E})$

# Computing primes

1.  $\mathcal{C}_m \wedge \mathcal{F} \not\models \perp$  — **tautology**
2.  $\mathcal{C}_m \wedge \mathcal{F} \models \mathcal{E} \Leftrightarrow \mathcal{C}_m \models (\mathcal{F} \rightarrow \mathcal{E})$



$\mathcal{C}_m$  is a **prime implicant** of  $\mathcal{F} \rightarrow \mathcal{E}$

## Computing one minimal explanation

- **One** subset-minimal explanation:

**Input:**  $\mathcal{F}$  under  $\mathcal{M}$ , initial cube  $\mathcal{C}$ , prediction  $\mathcal{E}$

**Output:** **Subset-minimal** explanation  $\mathcal{C}_m$

```
begin
    for  $I \in \mathcal{C}$  :
        if Entails( $\mathcal{C} \setminus \{I\}, \mathcal{F} \rightarrow \mathcal{E}$ ) :
             $\mathcal{C} \leftarrow \mathcal{C} \setminus \{I\}$ 
    return  $\mathcal{C}$ 
end
```

# Computing one minimal explanation

- **One** subset-minimal explanation:

**Input:**  $\mathcal{F}$  under  $\mathcal{M}$ , initial cube  $\mathcal{C}$ , prediction  $\mathcal{E}$

**Output:** **Subset-minimal** explanation  $\mathcal{C}_m$

begin

**for**  $I \in \mathcal{C}$  :

**if** Entails( $\mathcal{C} \setminus \{I\}, \mathcal{F} \rightarrow \mathcal{E}$ ) :

$\mathcal{C} \leftarrow \mathcal{C} \setminus \{I\}$

**return**  $\mathcal{C}$

end

- **One** cardinality-minimal explanation:

- Harder than computing subset-minimal explanation
- Exploit **implicit hitting set dualization**
- Details in earlier papers

# Outline

Successes & Pitfalls of ML

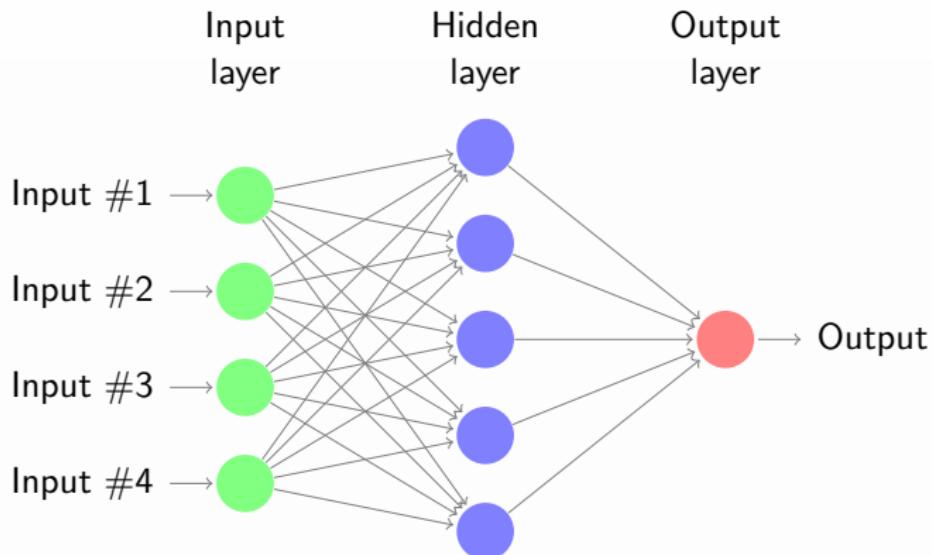
Explainable AI

Explanations with Abductive Reasoning

Encoding Neural Networks

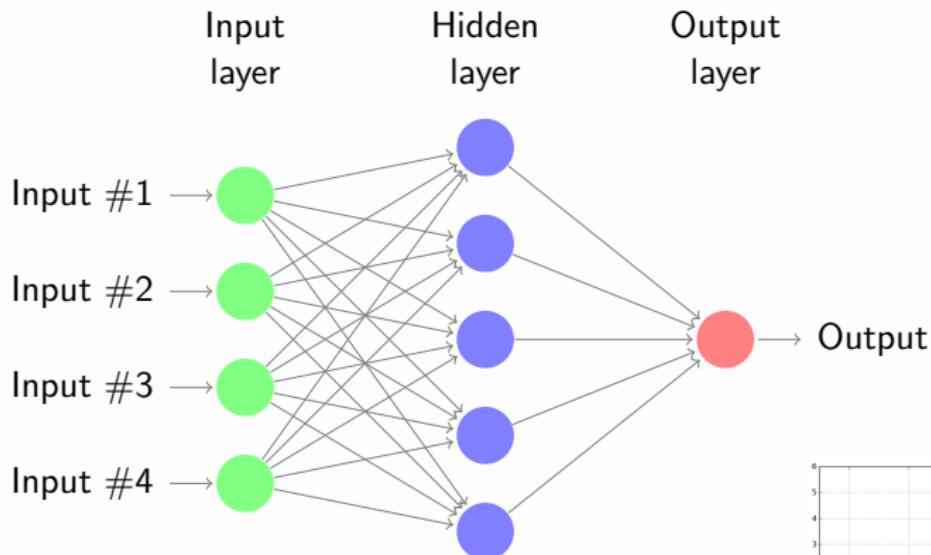
Results

## Encodings NNs

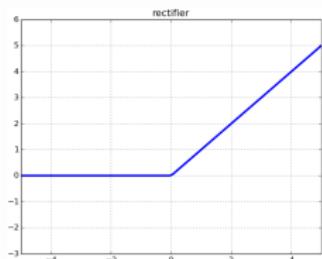


- Each layer (except first) viewed as a **block**
  - Compute  $\mathbf{x}'$  given input  $\mathbf{x}$ , weights matrix  $\mathbf{A}$ , and bias vector  $\mathbf{b}$
  - Compute output  $\mathbf{y}$  given  $\mathbf{x}'$  and activation function

# Encodings NNs



- Each layer (except first) viewed as a **block**
  - Compute  $\mathbf{x}'$  given input  $\mathbf{x}$ , weights matrix  $\mathbf{A}$ , and bias vector  $\mathbf{b}$
  - Compute output  $\mathbf{y}$  given  $\mathbf{x}'$  and activation function
- Each unit uses a **ReLU** activation function



# Encoding NNs using MILP

Computation for a NN ReLU **block**:

$$\mathbf{x}' = \mathbf{A} \cdot \mathbf{x} + \mathbf{b}$$

$$\mathbf{y} = \max(\mathbf{x}', \mathbf{0})$$

# Encoding NNs using MILP

Computation for a NN ReLU **block**:

$$\mathbf{x}' = \mathbf{A} \cdot \mathbf{x} + \mathbf{b}$$

$$\mathbf{y} = \max(\mathbf{x}', \mathbf{0})$$

Block **encoded** as follows:

[Fischetti&Jo, CJ'18]

$$\sum_{j=1}^n a_{i,j}x_j + b_i = y_i - s_i$$

$$z_i = 1 \rightarrow y_i \leq 0$$

$$z_i = 0 \rightarrow s_i \leq 0$$

$$y_i \geq 0, s_i \geq 0, z_i \in \{0, 1\}$$

- Simpler encodings not as effective

[Katz et al. CAV'17]

# Outline

Successes & Pitfalls of ML

Explainable AI

Explanations with Abductive Reasoning

Results

# Experimental setup

- Implementation in Python
  - Supports SMT solvers through PySMT
    - ▶ Yices2 used
  - Supports CPLEX 12.8.0

# Experimental setup

- Implementation in Python
    - Supports SMT solvers through PySMT
      - ▶ Yices2 used
    - Supports CPLEX 12.8.0
  - ReLU-based neural networks
    - One *hidden* layer with  $i \in \{10, 15, 20\}$  neurons
    - Pick NN that achieves *good* accuracy
- [Fischetti&Jo CJ'18]

# Experimental setup

- Implementation in Python
  - Supports **SMT** solvers through PySMT
    - ▶ **Yices2** used
  - Supports **CPLEX 12.8.0**
- **ReLU-based** neural networks [Fischetti&Jo CJ'18]
  - One *hidden* layer with  $i \in \{10, 15, 20\}$  neurons
  - Pick NN that achieves *good* accuracy
- Benchmarks selected from:
  - **UCI** Machine Learning Repository
  - **Penn** Machine Learning Benchmarks
  - **MNIST** Digits Database

# Experimental setup

- Implementation in Python
  - Supports **SMT** solvers through PySMT
    - ▶ **Yices2** used
  - Supports **CPLEX 12.8.0**
- **ReLU-based** neural networks [Fischetti&Jo CJ'18]
  - One *hidden* layer with  $i \in \{10, 15, 20\}$  neurons
  - Pick NN that achieves *good* accuracy
- Benchmarks selected from:
  - **UCI** Machine Learning Repository
  - **Penn** Machine Learning Benchmarks
  - **MNIST** Digits Database
- Machine configuration:
  - Intel Core i7 2.8GHz, 8GByte
  - Time limit — **1800s**
  - Memory limit — **4GByte**

## Sample of experimental results

Dataset		Minimal explanation			Minimum explanation		
		size	SMT (s)	MILP (s)	size	SMT (s)	MILP (s)
australian	(14)	<b>m</b>	1	0.03	0.05	—	—
		<b>a</b>	8.79	1.38	0.33	—	—
		<b>M</b>	14	17.00	1.43	—	—
backache	(32)	<b>m</b>	13	0.13	0.14	—	—
		<b>a</b>	19.28	5.08	0.85	—	—
		<b>M</b>	26	22.21	2.75	—	—
breast-cancer	(9)	<b>m</b>	3	0.02	0.04	3	0.02
		<b>a</b>	5.15	0.65	0.20	4.86	0.41
		<b>M</b>	9	6.11	0.41	9	1.81
cleve	(13)	<b>m</b>	4	0.05	0.07	4	—
		<b>a</b>	8.62	3.32	0.32	7.89	5.14
		<b>M</b>	13	60.74	0.60	13	39.06
hepatitis	(19)	<b>m</b>	6	0.02	0.04	4	0.01
		<b>a</b>	11.42	0.07	0.06	9.39	0.04
		<b>M</b>	19	0.26	0.20	19	2.89
voting	(16)	<b>m</b>	3	0.01	0.02	3	0.02
		<b>a</b>	4.56	0.04	0.13	3.46	0.25
		<b>M</b>	11	0.10	0.37	11	1.77
spect	(22)	<b>m</b>	3	0.02	0.02	3	0.04
		<b>a</b>	7.31	0.13	0.07	6.44	0.67
		<b>M</b>	20	0.88	0.29	20	10.73

## Sample of experimental results

Dataset		Minimal explanation			Minimum explanation		
		size	SMT (s)	MILP (s)	size	SMT (s)	MILP (s)
australian	(14)	m	1	0.03	0.05	—	—
		a	8.79	1.38	0.33	—	—
		M	14	17.00	1.43	—	—
backache	(32)	m	13	0.13	0.14	—	—
		a	19.28	5.08	0.85	—	—
		M	26	22.21	2.75	—	—
breast-cancer	(9)	m	3	0.02	0.04	3	0.02
		a	5.15	0.65	0.20	4.86	0.41
		M	9	6.11	0.41	9	1.81
cleve	(13)	m	4	0.05	0.07	4	—
		a	8.62	3.32	0.32	7.89	—
		M	13	60.74	0.60	13	39.06
hepatitis	(19)	m	6	0.02	0.04	4	0.01
		a	11.42	0.07	0.06	9.39	0.04
		M	19	0.26	0.20	19	2.89
voting	(16)	m	3	0.01	0.02	3	0.01
		a	4.56	0.04	0.13	3.46	0.02
		M	11	0.10	0.37	11	0.25
spect	(22)	m	3	0.02	0.02	3	0.02
		a	7.31	0.13	0.07	6.44	0.04
		M	20	0.88	0.29	20	0.67

## Sample of experimental results

Dataset		Minimal explanation			Minimum explanation		
		size	SMT (s)	MILP (s)	size	SMT (s)	MILP (s)
australian	(14)	m	1	0.03	0.05	—	—
		a	8.79	1.38	0.33	—	—
		M	14	17.00	1.43	—	—
backache	(32)	m	13	0.13	0.14	—	—
		a	19.28	5.08	0.85	—	—
		M	26	22.21	2.75	—	—
breast-cancer	(9)	m	3	0.02	0.04	3	0.02
		a	5.15	0.65	0.20	4.86	0.41
		M	9	6.11	0.41	9	24.80
cleve	(13)	m	4	0.05	0.07	4	—
		a	8.62	3.32	0.32	7.89	—
		M	13	60.74	0.60	13	39.06
hepatitis	(19)	m	6	0.02	0.04	4	0.01
		a	11.42	0.07	0.06	9.39	0.04
		M	19	0.26	0.20	19	2.89
voting	(16)	m	3	0.01	0.02	3	0.01
		a	4.56	0.04	0.13	3.46	0.02
		M	11	0.10	0.37	11	0.25
spect	(22)	m	3	0.02	0.02	3	0.02
		a	7.31	0.13	0.07	6.44	0.04
		M	20	0.88	0.29	20	0.67

## Sample of experimental results

Dataset		Minimal explanation			Minimum explanation		
		size	SMT (s)	MILP (s)	size	SMT (s)	MILP (s)
australian	(14)	m	1	0.03	0.05	—	—
		a	8.79	1.38	0.33	—	—
		M	14	17.00	1.43	—	—
backache	(32)	m	13	0.13	0.14	—	—
		a	19.28	5.08	0.85	—	—
		M	26	22.21	2.75	—	—
breast-cancer	(9)	m	3	0.02	0.04	3	0.02
		a	5.15	0.65	0.20	4.86	0.41
		M	9	6.11	0.41	9	1.81
cleve	(13)	m	4	0.05	0.07	4	—
		a	8.62	3.32	0.32	7.89	5.14
		M	13	60.74	0.60	13	39.06
hepatitis	(19)	m	6	0.02	0.04	4	0.04
		a	11.42	0.07	0.06	9.39	2.89
		M	19	0.26	0.20	19	22.23
voting	(16)	m	3	0.01	0.02	3	0.02
		a	4.56	0.04	0.13	3.46	0.25
		M	11	0.10	0.37	11	1.77
spect	(22)	m	3	0.02	0.02	3	0.04
		a	7.31	0.13	0.07	6.44	0.67
		M	20	0.88	0.29	20	10.73

# Comparing quality to compilation-based BNC

[Shih, Choi & Darwiche, IJCAI'18]

- “*Congressional Voting Records*” dataset

# Comparing quality to compilation-based BNC

[Shih, Choi & Darwiche, IJCAI'18]

- “*Congressional Voting Records*” dataset
- (0 1 0 1 1 1 0 0 0 0 0 0 1 1 0 1) — data sample (**16 features**)

# Comparing quality to compilation-based BNC

[Shih, Choi & Darwiche, IJCAI'18]

- “*Congressional Voting Records*” dataset
- (0 1 0 1 1 1 0 0 0 0 0 0 1 1 0 1) — data sample (**16 features**)

**smallest size** explanations computed by:

- ( 0 1 1 0 0 0 1 1 0 ) — **9 literals**
- ( 0 1 1 1 0 0 1 1 0 ) — **9 literals**

# Comparing quality to compilation-based BNC

[Shih, Choi & Darwiche, IJCAI'18]

- “*Congressional Voting Records*” dataset
- (0 1 0 1 1 1 0 0 0 0 0 0 1 1 0 1) — data sample (**16 features**)

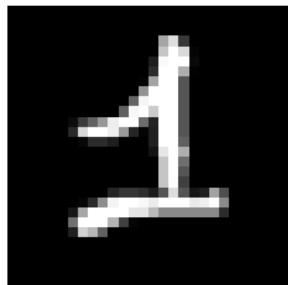
**smallest size** explanations computed by:

- ( 0 1 1 0 0 0 1 1 0 ) — **9 literals**
- ( 0 1 1 1 0 0 1 1 0 ) — **9 literals**

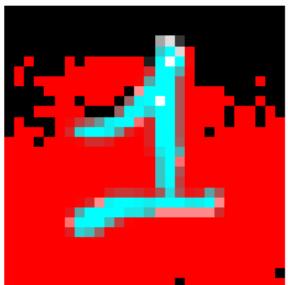
**subset-minimal** explanations computed by **our approach**:

- ( 1 0 0 0 ) — **4 literals**
- ( 1 0 0 ) — **3 literals**
- ( 0 1 0 0 0 ) — **5 literals**
- ( 0 1 0 0 1 ) — **5 literals**

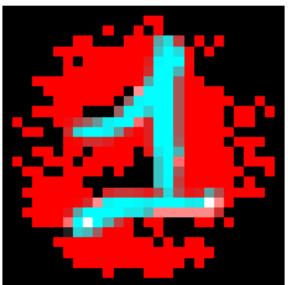
There are many explanations of different quality



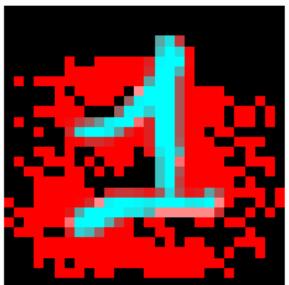
(a) digit 1



(b) simple expl.



(c) central pixels



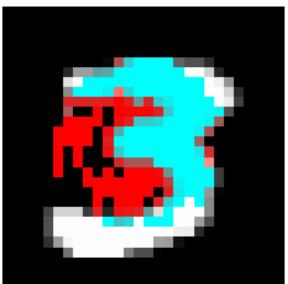
(d) light pixels



(a) digit 3



(b) simple expl.



(c) central pixels



(d) light pixels

# Outline

Successes & Pitfalls of ML

Explainable AI

Explanations with Abductive Reasoning

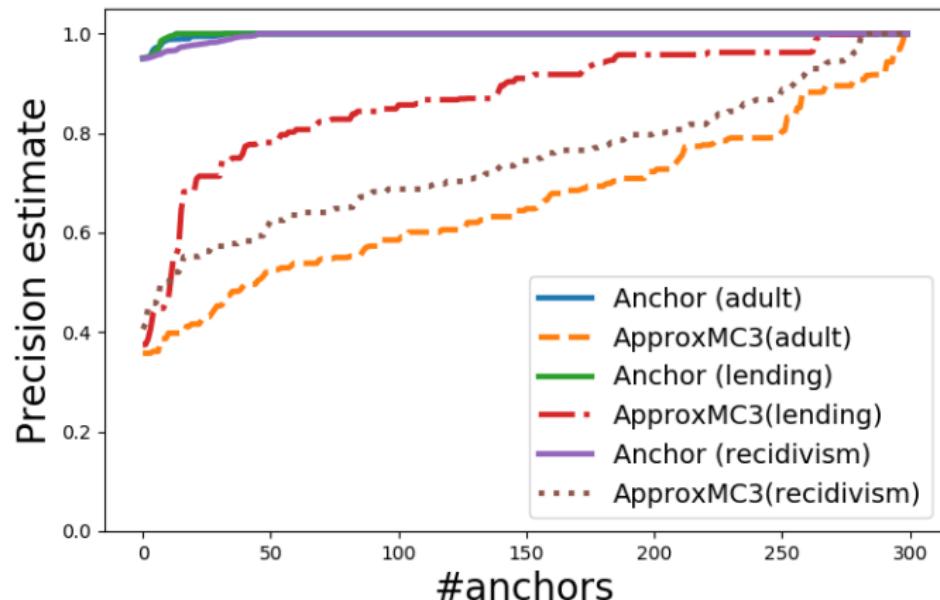
Results

Assessing Local Explanations – Recent Work

# Assessing precision with model counting

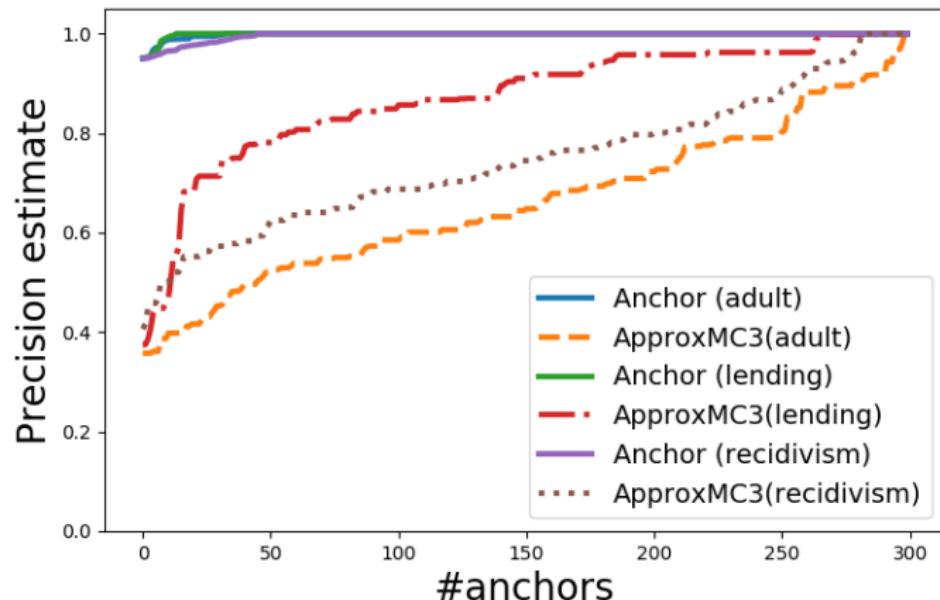
- Evaluated **Anchor** [Guerreiro et al., AAAI18]
  - Anchor more accurate than LIME
  - Anchor computes accuracy estimate for each explanation
- Represented ML model as **propositional formula**
  - E.g. **binarized NNs (BNNs)**
  - Use (approximate) model counter to assess precision of ML model on explanation (anchor) computed by **Anchor**

## Preliminary results



- Anchor often claims  $\approx 99\%$  precision

## Preliminary results



- Anchor often claims  $\approx 99\%$  precision; this cannot be confirmed

## Summary and future work

- **Principled** approach to XAI

## Summary and future work

- **Principled** approach to XAI
- Based on **abductive reasoning**

## Summary and future work

- **Principled** approach to XAI
- Based on **abductive reasoning**
- Applies a **reasoning engine**, e.g. **SMT** or **MILP**

## Summary and future work

- **Principled** approach to XAI
- Based on **abductive reasoning**
- Applies a **reasoning engine**, e.g. **SMT** or **MILP**
- Provides **minimality guarantees**

## Summary and future work

- **Principled** approach to XAI
- Based on **abductive reasoning**
- Applies a **reasoning engine**, e.g. **SMT** or **MILP**
- Provides **minimality guarantees**
- Tested on **ReLU-based** NNs
- First results on precision of Anchor's explanations

## Summary and future work

- **Principled** approach to XAI
- Based on **abductive reasoning**
- Applies a **reasoning engine**, e.g. **SMT** or **MILP**
- Provides **minimality guarantees**
- Tested on **ReLU-based** NNs
- First results on precision of Anchor's explanations
- **Other** ML models?

## Summary and future work

- **Principled** approach to XAI
- Based on **abductive reasoning**
- Applies a **reasoning engine**, e.g. **SMT** or **MILP**
- Provides **minimality guarantees**
- Tested on **ReLU-based** NNs
- First results on precision of Anchor's explanations
- **Other** ML models?
- Address scalability:
  - Better **encodings**?
  - More advanced **reasoners**?

## Summary and future work

- **Principled** approach to XAI
- Based on **abductive reasoning**
- Applies a **reasoning engine**, e.g. **SMT** or **MILP**
- Provides **minimality guarantees**
- Tested on **ReLU-based** NNs
- First results on precision of Anchor's explanations
- **Other** ML models?
- Address scalability:
  - Better **encodings**?
  - More advanced **reasoners**?
- Explanation **enumeration**? + **preferences**?

Questions?

## References to our work

- A. Ignatiev, N. Narodytska, J. Marques-Silva:  
Abduction-Based Explanations for Machine Learning Models.  
AAAI 2019
- N. Narodytska, A. Ignatiev, F. Pereira, J. Marques-Silva:  
Learning Optimal Decision Trees with SAT.  
IJCAI 2018: 1362-1368
- A. Ignatiev, F. Pereira, N. Narodytska, J. Marques-Silva:  
A SAT-Based Approach to Learn Explainable Decision Sets.  
IJCAR 2018: 627-645

## Additional references I

- M. T. Ribeiro, S. Singh, C. Guestrin:  
"Why Should I Trust You?": Explaining the Predictions of Any Classifier.  
KDD 2016: 1135-1144
- G. Katz, C. W. Barrett, D. L. Dill, K. Julian, M. J. Kochenderfer:  
Reluplex: An Efficient SMT Solver for Verifying Deep Neural Networks.  
CAV (1) 2017: 97-117
- M. T. Ribeiro, S. Singh, C. Guestrin:  
Anchors: High-Precision Model-Agnostic Explanations.  
AAAI 2018: 1527-1535
- A Shih, A. Choi, A. Darwiche:  
A Symbolic Approach to Explaining Bayesian Network Classifiers.  
IJCAI 2018: 5103-5111
- M. Fischetti, J. Jo:  
Deep neural networks and mixed integer linear optimization.  
Constraints 23(3): 296-309 (2018)

## Additional references II

- B. Goodman, S. R. Flaxman:  
European Union Regulations on Algorithmic Decision-Making and a “Right to Explanation”. AI Magazine 38(3): 50-57 (2017)
- A. M. Aung, Y. Fadila, R. Gondokaryono, L. Gonzalez:  
Building Robust Deep Neural Networks for Road Sign Detection.  
CoRR abs/1712.09327 (2017)
- K. Eykholt, I. Evtimov, E. Fernandes, B. Li, A. Rahmati, C. Xiao, A. Prakash, T. Kohno, D. Song:  
Robust Physical-World Attacks on Deep Learning Visual Classification.  
CVPR 2018: 1625-1634
- A. Madry, L. Schmidt:  
A Brief Introduction to Adversarial Examples.  
[http://gradientscience.org/intro\\_adversarial/](http://gradientscience.org/intro_adversarial/), 2018
- M. P. Kumar:  
Tutorial: Neural Network Verification.  
VMCAI 2019 Winter School  
<http://mpawankumar.info/tutorials/vmcai2019/>, 2019